

# CSCE 633: Machine Learning

## Lecture 20: More Boosting!

Texas A&M University

10-9-19

# Last Time

- Decision Trees
- Random Forest
- Boosting

# Goals of this lecture

- Reminder: Exam 1 - Monday, October 14 in class
- CLOSED BOOK, CLOSED NOTES - Starts right at 1:50. DO NOT BE LATE! YOU WILL NOT GET EXTRA TIME!
- Boosting

## Boosting: Formulation

Consider  $y \in \{-1, +1\}$  instead of  $\{0, 1\}$

Can calculate mean error in a classifier  $f$  as:

$$\bar{err} = \frac{1}{n} \sum_{i=1}^n \mathbb{I}(y_i \neq f(x_i))$$

Some definitions:

- A **weak classifier** is one that is just slightly better than random guessing.
- we define  $m = 1, 2, \dots, M$  as a sequence of weak classifier models
- You may see formulations of a strong classifier as  $G$  or  $H$  - usually capital to denote the strong classifier from weak
- You may see  $\alpha$  weights instead of  $w$  weights

## Boosting: Formulation

$$f(x) = \beta_0 + \sum_{b=1}^B \beta_b \phi_b(x)$$

Or - in other notation

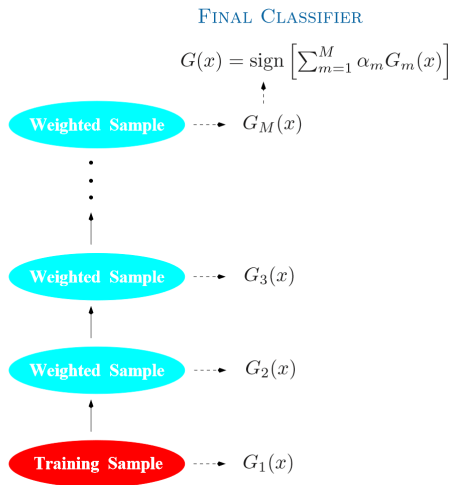
$$f(x) = w_0 + \sum_{m=1}^M w_m \phi_m(x)$$

Or - in other notation

$$G(x) = \text{sign}\left(\sum_{m=1}^M \alpha_m G_m(x)\right)$$

where  $\alpha_m$  are the weights on each weak classifier  $G_m$  and the binary classification is just the sum (regression would be mapped to a  $[-1, 1]$  interval in the same setting (called Real AdaBoost)

# AdaBoost



## Boosting: Minimizing Loss

Consider  $y \in \{-1, +1\}$  instead of  $\{0, 1\}$

Can calculate mean error in a classifier  $f$  as:

$$\bar{err} = \frac{1}{n} \sum_{i=1}^n \mathbb{I}(y_i \neq f(x_i))$$

Boosting aims to minimize loss as:

$$\min_f \sum_{i=1}^n L(y_i, f(x_i))$$

Where  $L(y_i, \hat{y}_i)$  is some loss function. What does this look like visually for binary classification?

## Boosting: Minimizing Loss

Consider  $y \in \{-1, +1\}$  instead of  $\{0, 1\}$

Binary 0 – 1 loss is not differentiable - so we need convex approximations.

Squared Error Loss

$$f^*(x) = \operatorname{argmin}_{f(x)} \mathbb{E}_{y|x}[(Y - f(x))^2] = \mathbb{E}[Y|x]$$

Cannot compute this because it requires  $p(y|x)$  to be known. This is commonly known as the population minimizer. Other loss functions with boosting, then, try to approximate this probability.



## Boosting: Other Loss Functions

Consider  $y \in \{-1, +1\}$  instead of  $\{0, 1\}$

Binary 0 – 1 loss is not differentiable - so we need convex approximations.

Squared Error Loss

$$f^*(x) = \operatorname{argmin}_{f(x)} \mathbb{E}_{y|x}[(Y - f(x))^2] = \mathbb{E}[Y|x]$$

Log Loss:

$$f^*(x) = \frac{1}{2} \log \frac{p(\tilde{y} = 1|x)}{p(\tilde{y} = -1|x)}$$

Exponential Loss:

$$L(\tilde{y}, f) = \exp(-\tilde{y}f)$$

Which has the same optimal estimate as log-loss it turns out!

## Boosting: Arriving at the optimal

Initially:

$$f_0(x) = \operatorname{argmin}_{\gamma} \sum_{i=1}^n L(y_i, f(x_i; \gamma))$$

For squared error, can start with  $f_0(x) = \bar{y}$  Then:

$$f_0(x) = \frac{1}{2} \log \frac{\hat{\pi}}{1 - \hat{\pi}}$$

where:

$$\hat{\pi} = \frac{1}{n} \sum_{i=1}^n \mathbb{I}(y_i = 1)$$

## Boosting: Iterating

Initially:

$$f_0(x) = \operatorname{argmin}_{\gamma} \sum_{i=1}^n L(y_i, f(x_i; \gamma))$$

For squared error, can start with  $f_0(x) = \bar{y}$  Then:

$$f_0(x) = \frac{1}{2} \log \frac{\hat{\pi}}{1 - \hat{\pi}}$$

Iterating:

$$f_m(x) = f_{m-1}(x) + \nu \beta_m \phi(x; \gamma_m)$$


Where  $0 < \nu \leq 1$  is a shrinkage parameter to make sure you learn slowly slowly!

Early Stopping: If accuracy (loss) does not improve on a validation set, or AIC, BIC, etc.

# Types of Loss

Name	Loss	Derivative	$f^*$	Algorithm
Squared error	$\frac{1}{2}(y_i - f(\mathbf{x}_i))^2$	$y_i - f(\mathbf{x}_i)$	$\mathbb{E}[y \mathbf{x}_i]$	L2Boosting
Absolute error	$ y_i - f(\mathbf{x}_i) $	$\text{sgn}(y_i - f(\mathbf{x}_i))$	$\text{median}(y \mathbf{x}_i)$	Gradient boosting
Exponential loss	$\exp(-\tilde{y}_i f(\mathbf{x}_i))$	$-\tilde{y}_i \exp(-\tilde{y}_i f(\mathbf{x}_i))$	$\frac{1}{2} \log \frac{\pi_i}{1-\pi_i}$	AdaBoost
Logloss	$\log(1 + e^{-\tilde{y}_i f(\mathbf{x}_i)})$	$y_i - \pi_i$	$\frac{1}{2} \log \frac{\pi_i}{1-\pi_i}$	LogitBoost

## Types of Loss



## AdaBoost M1

---

### Algorithm 10.1 *AdaBoost.M1*.

---

1. Initialize the observation weights  $w_i = 1/N$ ,  $i = 1, 2, \dots, N$ .
2. For  $m = 1$  to  $M$ :
  - (a) Fit a classifier  $G_m(x)$  to the training data using weights  $w_i$ .
  - (b) Compute
$$\text{err}_m = \frac{\sum_{i=1}^N w_i I(y_i \neq G_m(x_i))}{\sum_{i=1}^N w_i}.$$
  - (c) Compute  $\alpha_m = \log((1 - \text{err}_m)/\text{err}_m)$ .
  - (d) Set  $w_i \leftarrow w_i \cdot \exp[\alpha_m \cdot I(y_i \neq G_m(x_i))]$ ,  $i = 1, 2, \dots, N$ .
3. Output  $G(x) = \text{sign} \left[ \sum_{m=1}^M \alpha_m G_m(x) \right]$ .

## AdaBoost

(Freund and Schapire, 1997)

$$H \subseteq \{-1, +1\}^X.$$

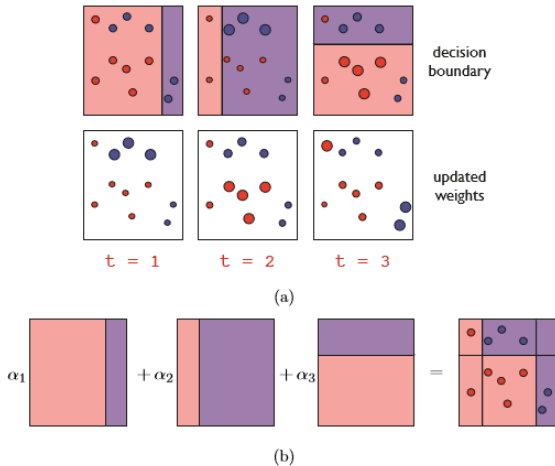
ADABOOST( $S = ((x_1, y_1), \dots, (x_m, y_m))$ )

```

1  for  $i \leftarrow 1$  to  $m$  do
2       $D_1(i) \leftarrow \frac{1}{m}$ 
3  for  $t \leftarrow 1$  to  $T$  do
4       $h_t \leftarrow$  base classifier in  $H$  with small error  $\epsilon_t = \Pr_{i \sim D_t} [h_t(x_i) \neq y_i]$ 
5       $\alpha_t \leftarrow \frac{1}{2} \log \frac{1 - \epsilon_t}{\epsilon_t}$ 
6       $Z_t \leftarrow 2[\epsilon_t(1 - \epsilon_t)]^{\frac{1}{2}}$   $\triangleright$  normalization factor
7      for  $i \leftarrow 1$  to  $m$  do
8           $D_{t+1}(i) \leftarrow \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t}$ 
9       $f_t \leftarrow \sum_{s=1}^t \alpha_s h_s$ 
10 return  $h = \text{sgn}(f_T)$ 

```

# AdaBoost M1



**Figure 7.2**

Example of AdaBoost with axis-aligned hyperplanes as base classifiers. (a) The top row shows decision boundaries at each boosting round. The bottom row shows how weights are updated at each round, with incorrectly (resp., correctly) points given increased (resp., decreased) weights. (b) Visualization of final classifier, constructed as a non-negative linear combination of base classifiers.



## AdaBoost Loss

Consider the exponential loss function:

$$L(y, f(x)) = e^{-yf(x)}$$

which means we want to solve:

$$(\beta_m, G_m) = \underset{\beta, G}{\operatorname{argmin}} \sum_{i=1}^N \exp[-y_i f_{m-1}(x_i) + \beta G(x_i)]$$

which is a representation of forward stagewise additive modeling

## Additive Models

---

**Algorithm 10.2** *Forward Stagewise Additive Modeling.*

---

1. Initialize  $f_0(x) = 0$ .

2. For  $m = 1$  to  $M$ :

(a) Compute

$$(\beta_m, \gamma_m) = \arg \min_{\beta, \gamma} \sum_{i=1}^N L(y_i, f_{m-1}(x_i) + \beta b(x_i; \gamma)).$$

(b) Set  $f_m(x) = f_{m-1}(x) + \beta_m b(x; \gamma_m)$ .

---

## AdaBoost Loss

$$(\beta_m, G_m) = \underset{\beta, G}{\operatorname{argmin}} \sum_{i=1}^N \exp[-y_i f_{m-1}(x_i) + \beta G(x_i)]$$

Can be re-written as:

$$(\beta_m, G_m) = \underset{\beta, G}{\operatorname{argmin}} \sum_{i=1}^N w_i^{(m)} \exp[-\beta y_i G(x_i)]$$

Which we can solve in two steps

## AdaBoost Loss

$$(\beta_m, G_m) = \operatorname{argmin}_{\beta, G} \sum_{i=1}^N w_i^{(m)} \exp[-\beta y_i G(x_i)]$$

For  $\beta > 0$ :

$$G_m = \operatorname{argmin}_G \sum_{i=1}^N w_i^{(m)} I(y_i \neq G(x_i))$$

Which minimizes misclassifications

## AdaBoost Loss

For  $\beta > 0$ :

$$G_m = \operatorname{argmin}_G \sum_{i=1}^N w_i^{(m)} I(y_i \neq G(x_i))$$

in other words

$$\begin{aligned} & e^{-\beta} \sum_{y_i = G(x_i)} w_i^{(m)} + e^{\beta} \sum_{y_i \neq G(x_i)} w_i^{(m)} \\ & (e^{\beta} - e^{-\beta}) \sum_{i=1}^N w_i^{(m)} I(y_i \neq G(x_i)) + e^{\beta} \sum_{i=1}^N w_i^{(m)} \end{aligned}$$

## AdaBoost Loss

Plugging in:

$$(\beta_m, G_m) = \underset{\beta, G}{\operatorname{argmin}} \sum_{i=1}^N \exp[-y_i f_{m-1}(x_i) + \beta G(x_i)]$$

yields:

$$\beta_m = \frac{1}{2} \log \frac{1 - \operatorname{err}_m}{\operatorname{err}_m}$$

,  
where  $\operatorname{err}_m$  is:

$$\operatorname{err}_m = \frac{\sum_{i=1}^N w_i^{(m)} I(y_i \neq G(x_i))}{\sum_{i=1}^N w_i^{(m)}}$$

## AdaBoost Loss

and then we update our approximate  $f$  as

$f_m(x) = f_{m-1}(x) + \beta_m G_m(x)$ , which updates weights as  
 $w_i^{(m+1)} = w_i^{(m)} e^{-\beta_m y_i G_m(x_i)}$

and since

$$-y_i G_m(x_i) = 2 * I(y_i \neq G_m(x_i)) - 1$$

$$w_i^{(m+1)} = w_i^{(m)} e^{\alpha_m I(y_i \neq G_m(x_i))} e_m^\beta$$

## Why does this work?

$$f^*(x) = \operatorname{argmin}_{f(x)} E_{Y|x}(e^{-Yf(x)}) = \frac{1}{2} \log \frac{\Pr(Y = 1|x)}{\Pr(Y = -1|x)}$$

or


$$\Pr(Y = 1|x) = \frac{1}{1 + e^{-2f^*(x)}}$$



# AdaBoost = Coordinate Descent

- **Objective Function:** convex and differentiable.

$$F(\bar{\alpha}) = \frac{1}{m} \sum_{i=1}^m e^{-y_i f(x_i)} = \frac{1}{m} \sum_{i=1}^m e^{-y_i \sum_{j=1}^N \bar{\alpha}_j h_j(x_i)} .$$



## Coordinate Descent

- **Direction:** unit vector  $\mathbf{e}_k$  with best directional derivative:

$$F'(\bar{\alpha}_{t-1}, \mathbf{e}_k) = \lim_{\eta \rightarrow 0} \frac{F(\bar{\alpha}_{t-1} + \eta \mathbf{e}_k) - F(\bar{\alpha}_{t-1})}{\eta}.$$

- Since  $F(\bar{\alpha}_{t-1} + \eta \mathbf{e}_k) = \sum_{i=1}^m e^{-y_i \sum_{j=1}^N \bar{\alpha}_{t-1,j} h_j(x_i) - \eta y_i h_k(x_i)}$ ,

$$\begin{aligned} F'(\bar{\alpha}_{t-1}, \mathbf{e}_k) &= -\frac{1}{m} \sum_{i=1}^m y_i h_k(x_i) e^{-y_i \sum_{j=1}^N \bar{\alpha}_{t-1,j} h_j(x_i)} \\ &= -\frac{1}{m} \sum_{i=1}^m y_i h_k(x_i) \bar{D}_t(i) \bar{Z}_t \\ &= -\left[ \sum_{i=1}^m \bar{D}_t(i) 1_{y_i h_k(x_i)=+1} - \sum_{i=1}^m \bar{D}_t(i) 1_{y_i h_k(x_i)=-1} \right] \frac{\bar{Z}_t}{m} \\ &= -\left[ (1 - \bar{\epsilon}_{t,k}) - \bar{\epsilon}_{t,k} \right] \frac{\bar{Z}_t}{m} = \boxed{2\bar{\epsilon}_{t,k} - 1} \frac{\bar{Z}_t}{m}. \end{aligned}$$

Thus, direction corresponding to base classifier with smallest error.


## Coordinate Descent

- **Step size:**  $\eta$  chosen to minimize  $F(\bar{\alpha}_{t-1} + \eta \mathbf{e}_k)$ ;

$$\begin{aligned}\frac{dF(\bar{\alpha}_{t-1} + \eta \mathbf{e}_k)}{d\eta} = 0 &\Leftrightarrow - \sum_{i=1}^m y_i h_k(x_i) e^{-y_i \sum_{j=1}^N \bar{\alpha}_{t-1,j} h_j(x_i)} e^{-\eta y_i h_k(x_i)} = 0 \\&\Leftrightarrow - \sum_{i=1}^m y_i h_k(x_i) \bar{D}_t(i) \bar{Z}_t e^{-\eta y_i h_k(x_i)} = 0 \\&\Leftrightarrow - \sum_{i=1}^m y_i h_k(x_i) \bar{D}_t(i) e^{-\eta y_i h_k(x_i)} = 0 \\&\Leftrightarrow - [(1 - \bar{\epsilon}_{t,k}) e^{-\eta} - \bar{\epsilon}_{t,k} e^{\eta}] = 0 \\&\Leftrightarrow \eta = \frac{1}{2} \log \frac{1 - \bar{\epsilon}_{t,k}}{\bar{\epsilon}_{t,k}}.\end{aligned}$$

Thus, step size matches base classifier weight of AdaBoost.


## Alternative Loss Functions



## Types of Loss

Name	Loss	Derivative	$f^*$	Algorithm
Squared error	$\frac{1}{2}(y_i - f(\mathbf{x}_i))^2$	$y_i - f(\mathbf{x}_i)$	$\mathbb{E}[y \mathbf{x}_i]$	L2Boosting
Absolute error	$ y_i - f(\mathbf{x}_i) $	$\text{sgn}(y_i - f(\mathbf{x}_i))$	$\text{median}(y \mathbf{x}_i)$	Gradient boosting
Exponential loss	$\exp(-\tilde{y}_i f(\mathbf{x}_i))$	$-\tilde{y}_i \exp(-\tilde{y}_i f(\mathbf{x}_i))$	$\frac{1}{2} \log \frac{\pi_i}{1-\pi_i}$	AdaBoost
Logloss	$\log(1 + e^{-\tilde{y}_i f(\mathbf{x}_i)})$	$y_i - \pi_i$	$\frac{1}{2} \log \frac{\pi_i}{1-\pi_i}$	LogitBoost

## Types of Loss



## Logit Boost

- Adaboost with exponential loss puts a lot of weight on misclassified examples.
- Additionally - it is hard to interpret probabilities from  $f(x)$
- If we use log-loss instead - mistakes are only punished linearly
- Can also generalize to multiple classes

$$p(y = 1|x) = \frac{e^{f(x)}}{e^{-f(x)} + e^{f(x)}} = \frac{1}{1 + e^{-2f(x)}}$$

$$L_m(\phi) = \sum_{i=1}^n \log(1 + \exp(-2\tilde{y}_i(f_{m-1}(x) + \phi(x_i))))$$

# Logit Boost: Algorithm

---

**Algorithm 16.3:** LogitBoost, for binary classification with log-loss

---

```
1  $w_i = 1/N, \pi_i = 1/2;$   
2 for  $m = 1 : M$  do  
3   Compute the working response  $z_i = \frac{y_i^* - \pi_i}{\pi_i(1 - \pi_i)};$   
4   Compute the weights  $w_i = \pi_i(1 - \pi_i);$   
5    $\phi_m = \operatorname{argmin}_{\phi} \sum_{i=1}^N w_i(z_i - \phi(\mathbf{x}_i))^2;$   
6   Update  $f(\mathbf{x}) \leftarrow f(\mathbf{x}) + \frac{1}{2}\phi_m(\mathbf{x});$   
7   Compute  $\pi_i = 1/(1 + \exp(-2f(\mathbf{x}_i)));$   
8 Return  $f(\mathbf{x}) = \operatorname{sgn} \left[ \sum_{m=1}^M \phi_m(\mathbf{x}) \right];$ 
```

---



## Gradient Descent Boosting

- Rather than rebuild the method per loss function, can we generalize?
- Imagine we want to minimize  $\hat{f} = \operatorname{argmin}_f L(f)$  where the  $f$  are the parameters of a model
- Then at step  $m$  let  $g_m$  be the gradient of  $L(f)$  at step  $f = f_{m-1}$

$$g_{im} = \left[ \frac{\partial L(y_i, f(x_i))}{\partial f(x_i)} \right]$$

# Types of Loss

Name	Loss	Derivative	$f^*$	Algorithm
Squared error	$\frac{1}{2}(y_i - f(\mathbf{x}_i))^2$	$y_i - f(\mathbf{x}_i)$	$\mathbb{E}[y \mathbf{x}_i]$	L2Boosting
Absolute error	$ y_i - f(\mathbf{x}_i) $	$\text{sgn}(y_i - f(\mathbf{x}_i))$	$\text{median}(y \mathbf{x}_i)$	Gradient boosting
Exponential loss	$\exp(-\tilde{y}_i f(\mathbf{x}_i))$	$-\tilde{y}_i \exp(-\tilde{y}_i f(\mathbf{x}_i))$	$\frac{1}{2} \log \frac{\pi_i}{1-\pi_i}$	AdaBoost
Logloss	$\log(1 + e^{-\tilde{y}_i f(\mathbf{x}_i)})$	$y_i - \pi_i$	$\frac{1}{2} \log \frac{\pi_i}{1-\pi_i}$	LogitBoost

## Functional Gradient Descent

$$g_{im} = \left[ \frac{\partial L(y_i, f(x_i))}{\partial f(x_i)} \right]$$

$$f_m = f_{m-1} - \rho_m g_m$$

where  $\rho_m$  is the step length and

$$\rho = \operatorname{argmin}_{\rho} L(f_{m-1} - \rho g_m)$$

But this does not generalize, only optimizes  $f$  for a fixed  $n$ . so we have to fit weak learners to approximate the negative gradient signal

$$\gamma_m = \operatorname{argmin}_{\gamma} \sum_{i=1}^n (-g_{im} - \phi(x_i; \gamma))^2$$

# Gradient Descent Boosting: Algorithm


---

**Algorithm 16.4:** Gradient boosting

---

- 1 Initialize  $f_0(\mathbf{x}) = \operatorname{argmin}_{\gamma} \sum_{i=1}^N L(y_i, \phi(\mathbf{x}_i; \gamma))$ ;
  - 2 **for**  $m = 1 : M$  **do**
  - 3     Compute the gradient residual using  $r_{im} = - \left[ \frac{\partial L(y_i, f(\mathbf{x}_i))}{\partial f(\mathbf{x}_i)} \right]_{f(\mathbf{x}_i) = f_{m-1}(\mathbf{x}_i)}$ ;
  - 4     Use the weak learner to compute  $\gamma_m$  which minimizes  $\sum_{i=1}^N (r_{im} - \phi(\mathbf{x}_i; \gamma_m))^2$ ;
  - 5     Update  $f_m(\mathbf{x}) = f_{m-1}(\mathbf{x}) + \nu \phi(\mathbf{x}; \gamma_m)$ ;
  - 6 **Return**  $f(\mathbf{x}) = f_M(\mathbf{x})$
-

## Boosting Comparisons



# Takeaways and Next Time

- Boosting
- Next Time: Discussion - Boosting and Exam Review
- Reminder: Exam 1 - Monday, October 14