

CSCE 633: Machine Learning

Lecture 26: Neural Networks

Texas A&M University

10-23-19

Last Time

- Clustering
- PROJECT PROPOSALS! GET WORKING :)

Goals of this lecture

- Introduction to Neural Networks

Neural networks: Original motivation

Inspiration from the brain

- Brain is a powerful information processing device
- Composed of a large number of processing units (neurons)
- Neurons operating *in parallel* → large *connectivity*
- Neural networks as a paradigm for parallel processing

Neural networks: Original motivation


Parallel computing architectures

- *Single Instruction Multiple Data (SIMD) machines*
 - All processors execute the same instruction but on different pieces of data
- *Multiple Instruction Multiple Data (MIMD) machines*
 - Different processors may execute different instructions on different data
- *Neural Instruction Multiple Data (NIMD) machines*
 - processors with small amount of local memory where some **parameters** can be stored
 - each processor implements a fixed function with different parameters
 - a little more complex than SIMD, but not as complex as MIMD

Perceptron


- Artificial Neural Networks have ability to do great things with vision, speech, learning
- Human brains are quite different from computers so how do we model?
- Neurons!
- Simple processing based upon activation, high levels of connectivity

Perceptron




- Each input has an associated weight (synaptic weight)
- $y = s \left(\sum_{j=1}^D w_d x_d + w_0 \right)$ where w_0 intercept makes the model more general - modeled as the weight coming from an extra bias unit (x_0) which is always $+1$.
- So, what is the learning procedure here?
- $y = s \left(\sum_{j=1}^D w_d x_d + w_0 \right)$ defines a hyperplane - so we can create a linear discriminant function to make decisions on classes.
- Unlike SVM - we can also get posterior probability using sigmoid as the output (like logistic regression)

Perceptron: Basic processing unit




- Inputs $x_d \in \mathbb{R}$, $d = 1, \dots, D$
 - might come from the environment
 - might be the output of other perceptrons
- Associated with a connection weight $w_d \in \mathbb{R}$, $d = 1, \dots, D$
- Output is some function of the linear combination of inputs
 - $y = s \left(\sum_{j=1}^D w_d x_d + w_0 \right) = s(\mathbf{w}^T \mathbf{x})$
where $s(\alpha) = 1$, if $\alpha > 0$, $s(\alpha) = 0$, otherwise
e.g. sigmoid activation: $s(\mathbf{x}, \mathbf{w}) = \frac{1}{1 + \exp(-\mathbf{w}^T \mathbf{x})}$
- can be used for classification, i.e. choose C_1 , if $s(\alpha) > 0$

Perceptron: Multiple classes



- For K classes, create K perceptrons
- Choose class C_i if $y_i = \max_k y_k$
- If we need probabilities, $o_i = w_i^T x$ which yields $y_i = \frac{\exp o_i}{\sum_k \exp o_k}$ called the softmax values

Perceptron: Basic processing unit



- Multiclass: $K > 2$ outputs
 - $y_k = s\left(\sum_{d=1}^D w_{kd}x_d + w_{k0}\right) = s(\mathbf{w}_k^T \mathbf{x})$
where w_{kj} is the weight from input x_j to output y_k
e.g. $s(\mathbf{x}, \mathbf{w}_1, \dots, \mathbf{w}_K) = \frac{\exp(\mathbf{w}_k^T \mathbf{x})}{1 + \sum_{k=1}^K \exp(\mathbf{w}_k^T \mathbf{x})}$
 - 0/1 encoding for output vector
 - e.g. in a 4-class problem: if class=3, then $y = [0, 0, 1, 0]$

Perceptron: Training

Online training

- Cost-efficient (computationally and memory-wise)
- Nature of data can change over time
- Error function expressed in terms of individual samples
- Weight update performed after each instance is seen

Perceptron: Training

Online training

- Evaluation: cross-entropy function for 1 instance (\mathbf{x}_n, y_n)

$$\mathcal{E}(\mathbf{w}) = -y_n \log [\sigma(\mathbf{w}^T \mathbf{x}_n)] - (1 - y_n) \log [1 - \sigma(\mathbf{w}^T \mathbf{x}_n)]$$

$$\mathcal{E}(\mathbf{w}_1, \dots, \mathbf{w}_K) = - \sum_{k=1}^K y_{nk} \log p(y_{nk} = 1 | \mathbf{w}_1, \dots, \mathbf{w}_K)$$

- Optimization: gradient descent


$$\frac{\partial \mathcal{E}(\mathbf{w})}{\partial w_d} = (\sigma(\mathbf{w}^T \mathbf{x}_n) - y_n) x_{nd}$$

$$\frac{\partial \mathcal{E}(\mathbf{w})}{\partial w_{kd}} = (\sigma(\mathbf{w}^T \mathbf{x}_n) - y_{nk}) x_{nd}$$

Approximating linear functions

Example: Boolean AND

x_1	x_2	r
0	0	0
0	1	0
1	0	0
1	1	1



Example of a perceptron implementing AND

$$y = s(x_1 + x_2 - 1.5)$$


$$\mathbf{w} = [-1.5 \ 1 \ 1]^T$$

$$\mathbf{x} = [1 \ x_1 \ x_2]^T$$

Approximating linear functions

Example: Boolean OR

x_1	x_2	r
0	0	0
0	1	1
1	0	1
1	1	1



Example of a perceptron implementing OR

$$y = s(x_1 + x_2 - 0.5)$$

$$\mathbf{w} = [-0.5 \ 1 \ 1]^T$$

$$\mathbf{x} = [1 \ x_1 \ x_2]^T$$

Approximating linear functions

Example: Boolean NOT

x_1	r
0	1
1	0

Example of a perceptron implementing NOT

$y = ?$


$\mathbf{w} = ?$

$\mathbf{x} = [1 \ x_1]^T$

Approximating linear functions

Example: Boolean NOT

x_1	r
0	1
1	0



Example of a perceptron implementing OR

$$y = s(x_1 - 2)$$

$$\mathbf{w} = [1 \ -2]^T$$

$$\mathbf{x} = [1 \ x_1]^T$$

Approximating linear functions

Example: Boolean (NOT x_1) AND (NOT x_2)

x_1	x_2	r
0	0	1
0	1	0
1	0	0
1	1	0

Example of a perceptron implementing OR

$y = ?$


$\mathbf{w} = ?$

$\mathbf{x} = [1 \ x_1 \ x_2]^T$

Approximating linear functions

Example: Boolean (NOT x_1) AND (NOT x_2)

x_1	x_2	r
0	0	1
0	1	0
1	0	0
1	1	0



Example of a perceptron implementing OR

$$y = s(-20x_1 - 20x_2 + 10)$$


$$\mathbf{w} = [10 \ -20 \ -20]^T$$

$$\mathbf{x} = [1 \ x_1 \ x_2]^T$$

Approximating linear functions

Example: Boolean XOR

x_1	x_2	r
0	0	0
0	1	1
1	0	1
1	1	0



Example of a perceptron implementing OR

$y = ?$


$\mathbf{w} = ?$

$\mathbf{x} = [1 \ x_1 \ x_2]^T$

Approximating linear functions

Example: Boolean XOR

x_1	x_2	r
0	0	0
0	1	1
1	0	1
1	1	0




Not linearly separable


Need combination of more than one perceptrons → **multilayer perceptrons**

Multilayer Perceptron

- Type of feedforward neural network
- Can model non-linear associations
- “Multi-level combination” of many perceptrons



Multilayer Perceptron: Representation



$$\alpha_1^{(2)} = f(W_{11}^{(1)}x_1 + W_{12}^{(1)}x_2 + W_{13}^{(1)}x_3 + b_1^{(1)})$$

$$\alpha_2^{(2)} = f(W_{21}^{(1)}x_1 + W_{22}^{(1)}x_2 + W_{23}^{(1)}x_3 + b_2^{(1)})$$

$$\alpha_3^{(2)} = f(W_{31}^{(1)}x_1 + W_{32}^{(1)}x_2 + W_{33}^{(1)}x_3 + b_3^{(1)})$$

$$h_{W,b}(x) = \alpha_1^{(3)} = f(W_{11}^{(2)}\alpha_1^{(2)} + W_{12}^{(2)}\alpha_2^{(2)} + W_{13}^{(2)}\alpha_3^{(2)} + b_1^{(2)})$$

Terminology

$W_{ij}^{(l)}$: connection between unit j in layer l to unit i in layer $l + 1$


$\alpha_i^{(l)}$: activation of unit i in layer l

$b_i^{(l)}$: bias connected with unit i in layer $l + 1$

Forward propagation: The process of propagating the input to the output through the activation of inputs and hidden units to each node

Multilayer Perceptron: Representation

Matrix notation



$$\alpha^{(2)} = f(\mathbf{W}^{(1)}\mathbf{x} + \mathbf{b}^{(1)})$$

$$h_{W,b}(\mathbf{x}) = \alpha^{(3)} = f(\mathbf{W}^{(2)}\alpha^{(2)} + \mathbf{b}^{(2)})$$


$$\mathbf{W}^{(1)} = \begin{bmatrix} W_{11}^{(1)} & W_{12}^{(1)} & W_{13}^{(1)} \\ W_{21}^{(1)} & W_{22}^{(1)} & W_{23}^{(1)} \\ W_{31}^{(1)} & W_{32}^{(1)} & W_{33}^{(1)} \end{bmatrix}, \mathbf{b}^{(1)} = [b_1^{(1)} \ b_2^{(1)} \ b_3^{(1)}], \text{ etc.}$$

Multilayer Perceptron: Representation

Alternative architectures


2 hidden layers, multiple output units

e.g. medical diagnosis: different outputs might indicate presence or absence of different diseases



Multilayer Perceptron



Non-linear feature learning



Multilayer Perceptron: Approximating non-linear functions


Example: Boolean XOR with multilayer perceptrons

x_1	x_2	z_1	z_2	r
0	0	0	1	1
0	1	0	0	0
1	0	0	0	0
1	1	1	0	1



Multilayer Perceptron


Question: How many parameters does the following network have to learn?



- A) 20
- B) 26
- C) 6
- D) 12

Multilayer Perceptron

Question: How many parameters does the following network have to learn?



- A) 20
- B) 26
- C) 6
- D) 12

The correct answer is B

$[3 \times 4] + [4 \times 2] = 20$ weights, $4 + 2 = 6$ biases

B Mortazavi CSE

Backpropagation

Multilayer Perceptron: Representation

- **Input:** $\mathbf{x} \in \mathbb{R}^D$
- **Output:**
 $y \in \{0, 1\}$ or $y \in \{1, \dots, K\}$ (classification)
 $y \in \mathbb{R}$ or $y \in \mathbb{R}^K$ (regression)
- **Training data:** $\mathcal{D}^{train} = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)\}$
- **Model:** $h_{\mathbf{W}, \mathbf{b}}(\mathbf{x})$
represented through forward propagation (see previous slides)
- **Model parameters:** weights $\mathbf{W}^{(1)}, \dots, \mathbf{W}^{(L)}$ and biases $\mathbf{b}^{(1)}, \dots, \mathbf{b}^{(L)}$

Multilayer Perceptron: Evaluation criterion

$$J(\mathbf{W}, \mathbf{b}, \mathcal{D}^{train}) = \frac{1}{2} \|h_{\mathbf{W}, \mathbf{b}}(\mathbf{x}) - y\|_2^2 \text{ (regression)}$$

$$J(\mathbf{W}, \mathbf{b}, \mathcal{D}^{train}) = y \log h_{\mathbf{W}, \mathbf{b}}(\mathbf{x}) + (1 - y) \log(1 - h_{\mathbf{W}, \mathbf{b}}(\mathbf{x})) \text{ (classification)}$$

Backpropagation: Intuition

- Consider a multi-layer perceptron
- Each layer is a perceptron with weights from the prior layer
- So error at the end based upon weights at the beginning is $\frac{\partial E}{\partial w_{hj}}$
- So by the chain rule:

$$\frac{\partial E}{\partial w_{hj}} = \frac{\partial E}{\partial y_i} \frac{\partial y_i}{\partial z_h} \frac{\partial z_h}{\partial w_{hi}}$$

Backpropagation

Multilayer Perceptron: Evaluation criterion

Regression

$$J(\mathbf{W}, \mathbf{b}) = \frac{1}{N} \sum_{n=1}^M \frac{1}{2} \|h_{\mathbf{W}, \mathbf{b}}(\mathbf{x}_n) - y_n\|_2^2 + \frac{\lambda}{2} \sum_{l=1}^{L-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} (W_{ji}^{(l)})^2$$

Classification

$$J(\mathbf{W}, \mathbf{b}) = \frac{1}{N} \sum_{n=1}^M (y_n \log h_{\mathbf{W}, \mathbf{b}}(\mathbf{x}_n) + (1 - y_n) \log(1 - h_{\mathbf{W}, \mathbf{b}}(\mathbf{x}_n))) \\ + \frac{\lambda}{2} \sum_{l=1}^{L-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} (W_{ji}^{(l)})^2$$

We will perform **gradient descent**

Backpropagation

Gradient descent for regression

$$J(\mathbf{W}, \mathbf{b}) = \frac{1}{N} \sum_{n=1}^M \frac{1}{2} \|h_{\mathbf{W}, \mathbf{b}}(\mathbf{x}_n) - y_n\|_2^2 + \frac{\lambda}{2} \sum_{l=1}^{L-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} (W_{ji}^{(l)})^2$$

$$W_{ij}^{(l)} := W_{ij}^{(l)} - \alpha \frac{\partial J(\mathbf{W}, \mathbf{b})}{\partial W_{ij}^{(l)}}$$

$$b_i^{(l)} := b_i^{(l)} - \alpha \frac{\partial J(\mathbf{W}, \mathbf{b})}{\partial b_i^{(l)}}$$

Note: Initialize the parameters randomly \rightarrow **symmetry breaking**

Use **backpropagation** to compute partial derivatives $\frac{\partial J(\mathbf{W}, \mathbf{b})}{\partial W_{ij}^{(l)}}$ and $\frac{\partial J(\mathbf{W}, \mathbf{b})}{\partial b_i^{(l)}}$

Backpropagation

Intuition

- Given a training example (\mathbf{x}_n, y_n) , we run a "forward pass" to compute all the activations
- For each node i in layer l , we compute an **error term** $\delta_i^{(l)}$ that measures how much that node was "responsible" for any errors in the output
 - Output node: difference between activation and target value
 - Hidden nodes: weighted average of the error terms of the nodes from the previous layer (i.e. $l + 1$)

Backpropagation


Implementation

- Given a training example (\mathbf{x}_n, y_n) , we run a "forward pass" to compute all the activations
- For each node i in output layer L
 - $\delta_i^{(L)} = (y_n - \alpha_i^{(L)})f'(z_i^{(L)})$
- For each node i in layer $l = L - 1, L - 2, \dots, 2$
 - Hidden nodes: $\delta_i^{(l)} = \left(\sum_{j=1}^{s_{l+1}} W_{ji}^{(l)} \delta_j^{(l+1)} \right) f'(z_i^{(l)})$
- Compute the desired partial derivatives as:
$$\frac{\partial J(\mathbf{W}, \mathbf{b})}{\partial W_{ij}^{(l)}} = \alpha_j^{(l)} \delta_i^{(l+1)}$$
$$\frac{\partial J(\mathbf{W}, \mathbf{b})}{\partial b_i^{(l)}} = \delta_i^{(l+1)}$$

Determining number of layers and their sizes

Implementation


- The capacity of the network (i.e. the number of representable functions) increases as we increase the number of layers
- How to avoid overfitting?



Determining number of layers and their sizes

How to avoid overfitting

- Limit # layers and #hidden units per layers
- Early stopping: start with small weights and stop learning early
- Weight decay: penalize large weights (regularization)
- Noise: add noise to the weights
- Add constraints to the weights



The effects of regularization strength: Each neural network above has 20 hidden neurons, but changing the regularization strength makes its final decision regions smoother with a higher regularization. You can play with these examples in this [ConvNetsJS demo](http://cs.stanford.edu/people/karpathy/convnetjs/demo/classify2d.html).

5

<http://cs.stanford.edu/people/karpathy/convnetjs/demo/classify2d.html>

Determining number of layers and their sizes

How to chose the number of layers and nodes

- No general rule of thumb, this depends on:
 - Amount of training data available
 - Complexity of the function that is trying to be learned
 - Number of input and output nodes
- If data is linearly separable, you don't need any hidden layers at all
- Start with one layer and hidden nodes proportional to input size
- Gradually increase

Activation Function


Transforms the activation level of a node (weighted sum of inputs) to an output signal

- Sigmoid: $\sigma(x) = \frac{1}{1+e^{-x}}$
- Hyperbolic tangent: $s(x) = \tanh(x) = 2\sigma(2x) - 1$
- Rectified Linear Unit (ReLU): $f(x) = \max(0, x)$
- Leaky ReLU: $f(x) = (ax) \cdot \mathbb{I}(x < 0) + (x) \cdot \mathbb{I}(x \geq 0)$ (e.g. $a = 0.01$)

Activation Function

Sigmoid: $s(x) = \frac{1}{1+e^{-x}}$


- Transforms a real-valued number between 0 and 1
- Large negative numbers become 0 (not firing at all)
- Large positive numbers become 1 (fully-saturated firing)
- Used historically because of its nice interpretation
- **Saturates gradients:** The gradient at either extremes (0 or 1) is almost zero, “killing” the signal will flow
- **Non-zero centered output:** Can be problematic during training, since it can bias outputs toward being always positive or always negative, causing unnecessary oscillations during the optimization



Activation Function

Hyperbolic tangent: $s(x) = \tanh(x) = 2\sigma(2x) - 1$


- Scaled version of sigmoid
- Transforms a real-valued number between -1 and 1
- Saturates gradients: Similar to sigmoid
- Output is zero-centered, avoiding some oscillation issues



Activation Function

Rectified Linear Unit (ReLU): $f(x) = \max(0, x)$


- Activation simply thresholded at zero
- Very popular during the last years
- **Accelerates convergence** (e.g. a factor of 6, see bellow) compared to the sigmoid/tanh (due to its linear, non-saturating form)
- **Cheap implementation** by simply thresholding at zero
- Activation can “die”: a large gradient flowing through a ReLU neuron could cause the weights to update in such a way that the neuron will never activate on any datapoint again, proper adjustment of learning rate can mitigate that



Activation Function

Leaky ReLU: $f(x) = (ax) \cdot \mathbb{I}(x < 0) + (x) \cdot \mathbb{I}(x \geq 0)$

- Instead of the function being zero when $x < 0$, leaky ReLU will have a small negative slope (e.g. $a = 0.01$)
- Some successful results, but not always consistent



Takeaways and Next Time

- Perceptron
- Back propagation
- Activation
- Next Time: More Neural Networks